

# Gaussian Processes

---

If we want to do better, we need a better predictor

**What would be some desired properties?**

## Filling Values Using a Model

**Given a gap (i.e. one or more contiguous missing values), the model:**

- Should be able to make a prediction about the missing values
- It should take into account the values that are **before** the hole
- ...But also those that are **after** the hole

I.e. it should be able to interpolate **all the available data**

**Can you think of an ML model that can do that?**

## Filling Values Using a Model

**Given a gap (i.e. one or more contiguous missing values), the model:**

- Should be able to make a prediction about the missing values
- It should take into account the values that are **before** the hole
- ...But also those that are **after** the hole

I.e. it should be able to interpolate **all the available data**

**Can you think of an ML model that can to that?**

If you have trouble, that's normal: there aren't many!

- Most ML models **cannot** be used for filling (at least not easily)
- We will see a couple of examples now

# Filling Values Using a Density Estimator

Density estimation does take into account all data

...But it cannot easily provide predictions

To be fair, predictions **can be obtained** from a density estimator

- Given an estimator  $f(x, \theta)$  for  $P(x)$
- ...Then we can find the most likely value for  $x$  by solving:

$$\operatorname{argmax}_x f(x, \theta)$$

- This is called a **Maximum A Posteriori (MAP)**

In fact, it's what we naturally get out of most regressors/classifiers

# Filling Values Using a Density Estimator

Density estimation does take into account all data

...But it cannot easily provide predictions

To be fair, predictions **can be obtained** from a density estimator

- Given an estimator  $f(x, \theta)$  for  $P(x)$
- ...Then we can find the most likely value for  $x$  by solving:

$$\operatorname{argmax}_x f(x, \theta)$$

- This is called a **Maximum A Posteriori (MAP)**

In fact, it's what we naturally get out of most regressors/classifiers

...But with density estimators, computing the MAP can be **very expensive**

# Filling Values Using a Regressor

**We can train a (auto)-regressor to predict the next value in a series**

...Then we can fill holes by making a "rolling forecast"

- We make one prediction
- We incorporate the prediction in the input
- Then we make another prediction and so on

...But by doing so we only rely on past values

**In fact, the last prediction can be far from the next true value**


- In addition, every iteration of the process introduces approximation
- Formally, we get compound errors

This is partly unavoidable: predicting the far future is often difficult

# Gaussian Processes

One of the few viable ML models is given by **Gaussian Processes (GP)**

They are not easy to use, but can work well in the correct setting




- GPs define a probability distribution over an **index (i.e. input) variable**
- The distribution is based on **the available observation** and **a few assumptions**



# Gaussian Processes

One of the few viable ML models is given by **Gaussian Processes (GP)**




## Assumption 1 (intuitively)

- For every value of the index variable the distribution is Gaussian
- Therefore it can be described by a mean and a standard deviation

# Gaussian Processes

One of the few viable ML models is given by **Gaussian Processes (GP)**



## Assumption 2 (intuitively)

- The stdev depends on the distance between a point and the observations
- So it will be **low** when we are **close** to the observations, **high** when we are **far**

# Gaussian Processes

## Formally, things are a bit more complicated

A GP is a **stochastic process**, i.e. a collection of indexed random variables

- The index variable  $\mathbf{x}$  represents an input (and can be a vector/tuple)
- Each variable  $y_{\mathbf{x}}$  represents the output for input  $\mathbf{x}$
- The index is **continuous** and the collection is therefore **infinite**

You can think of  $y_{\mathbf{x}}$  as the value of a (stochastic) function for input  $\mathbf{x}$

## Some examples:

- $y_{\mathbf{x}}$  could be the rainfall rate at location  $\mathbf{x} = (\text{latitude}, \text{longitude})$
- $y_{\mathbf{x}}$  could be the traffic volume at time  $\mathbf{x}$

## Each $y_{\mathbf{x}}$ follows a Normal Distribution, but the variables are **correlated**

- Therefore every finite subset of  $y_{\mathbf{x}}$  variables
- ...Follows a **Multivariate Normal Distribution**

# Multivariate Normal Distribution

## Why the multivariate normal distribution?

- It works for many real world phenomena
- It has a (relatively) simple closed-form density function

# Multivariate Normal Distribution

## Why the multivariate normal distribution?

- It works for many real world phenomena
- It has a (relatively) simple closed-form density function

## In particular, the PDF for a MND is defined via:

- A (vector) mean  $\mu$
- A covariance matrix  $\Sigma$

By recentering we can assume  $\mu = 0$ , meaning that knowing  $\Sigma$  is enough

# Multivariate Normal Distribution

## Why the multivariate normal distribution?

- It works for many real world phenomena
- It has a (relatively) simple closed-form density function

## In particular, the PDF for a MND is defined via:

- A (vector) mean  $\mu$
- A covariance matrix  $\Sigma$


By recentering we can assume  $\mu = 0$ , meaning that **knowing  $\Sigma$  is enough**

## Therefore, if we know $\Sigma$ we can easily compute

- The joint density  $f(\bar{y}_{\bar{x}})$  for a set of observations
- The conditional density  $f(y_x \mid \bar{y}_{\bar{x}})$  of an observation  $y_x$ , given  $\bar{y}_{\bar{x}}$

# Why do We Care?

We need the conditional density to perform inference!



- The line and grey areas represent the conditional density  $f(y_x | \bar{y}_{\bar{x}})$  of  $y_x$
- ...Based on the available observations, i.e.  $\bar{y}_{\bar{x}}$

# Why do We Care

**We need the joint density to perform training!**

...Because in practice we **don't know  $\Sigma$**

- Therefore we'll assume that  $\Sigma$  is a **parameterized function  $\Sigma(\theta)$**
- ...And we can optimize the parameters  $\theta$  for maximum likelihood



# Why do We Care

**We need the joint density to perform training!**

...Because in practice we **don't know  $\Sigma$**

- Therefore we'll assume that  $\Sigma$  is a **parameterized function  $\Sigma(\theta)$**

- ...And we can optimize the parameters  $\theta$  for maximum likelihood

**Formally, given a set of training observations  $\bar{y}_{\bar{x}}$**

...We can calibrate the parameters by solving a problem in the form:

$$\arg \max_{\theta} f(\bar{y}_{\bar{x}})$$

- Here we are not using a product of probabilities over the training set

- ...Because the  $\mathbf{y}_x$  variables are correlated

- No worries: since we have the joint PDF, we use it directly

# Which Covariance Matrix?

## We still have one big problem

Say that we have a covariance matrix  $\Sigma$  for a set of observations  $\bar{y}_{\bar{x}}$

- Now we want to perform inference for an input value  $x$
- Formally: we want to compute  $f(y_x \mid \bar{y}_{\bar{x}})$

In principle, we can use the formula:

$$f(y_x \mid \bar{y}_{\bar{x}}) = \frac{f(y_x, \bar{y}_{\bar{x}})}{f(\bar{y}_{\bar{x}})}$$

- By using our  $\Sigma$ , we can easily compute  $f(\bar{y}_{\bar{x}})$
- ...But what about  $f(y_x, \bar{y}_{\bar{x}})$ ?

## Which Covariance Matrix?

Our  $\Sigma$  refers to our set of observed variables  $\bar{y}_{\bar{x}}$

Therefore, if we have  $n$  variables, our matrix will be  $n \times n$

$$\Sigma_{\bar{x}} = \begin{pmatrix} \sigma_{\bar{x}_1, \bar{x}_1} & \sigma_{\bar{x}_1, \bar{x}_2} & \cdots & \sigma_{\bar{x}_1, \bar{x}_n} \\ \sigma_{\bar{x}_2, \bar{x}_1} & \sigma_{\bar{x}_2, \bar{x}_2} & \cdots & \sigma_{\bar{x}_2, \bar{x}_n} \\ \vdots & \vdots & \ddots & \vdots \\ \sigma_{\bar{x}_n, \bar{x}_1} & \sigma_{\bar{x}_n, \bar{x}_2} & \cdots & \sigma_{\bar{x}_n, \bar{x}_n} \end{pmatrix}$$

- In every cell we have the covariance for variables  $\bar{y}_{\bar{x}_i}$  and  $\bar{y}_{\bar{x}_j}$
- With this matrix, we can compute  $f(\bar{y}_{\bar{x}})$

bar## Which Covariance Matrix?

**However,  $f(y_x, \bar{y}_{\bar{x}})$  refers to one more variable**

Meaning that it will be specified via an  $(n + 1) \times (n + 1)$  matrix!

$$\Sigma_{x, \bar{x}} = \begin{pmatrix} \sigma_{x,x} & \sigma_{x,\bar{x}_1} & \sigma_{x,\bar{x}_2} & \cdots & \sigma_{x,\hat{x}_2} \\ \sigma_{\bar{x}_1,x} & \sigma_{\bar{x}_1,\bar{x}_1} & \sigma_{\bar{x}_1,\bar{x}_2} & \cdots & \sigma_{\hat{x}_1,\hat{x}_n} \\ \sigma_{\bar{x}_2,x} & \sigma_{\bar{x}_2,\bar{x}_1} & \sigma_{\bar{x}_2,\bar{x}_2} & \cdots & \sigma_{\bar{x}_2,\bar{x}_n} \\ \vdots & \vdots & \ddots & \vdots & \\ \sigma_{\bar{x}_n,x} & \sigma_{\bar{x}_n,\bar{x}_1} & \sigma_{\bar{x}_n,\bar{x}_2} & \cdots & \sigma_{\bar{x}_n,\bar{x}_2} \end{pmatrix}$$

- Assuming that  $\bar{y}_{\bar{x}}$  are the training observations
- ...We could define  $\sigma_{\bar{x}_1,\bar{x}_2}$  at training time

**But how do we define the new covariances, i.e. those related to  $y_x$ ?**

# Defining the Covariance Matrix

We assume that covariance can be **built from a set of inputs**

Let  $\mathbf{x}$  refer now to a vector of values for the index variable

- Given two variables  $y_{x_i}$  and  $y_{x_j}$
- We specify their covariance via parameterized **kernel function**  $K_{\theta}(x_i, x_j)$
- $K$  typically depends on the distance between input values

**Given any finite set of variables  $\{y_{x_1}, \dots, y_{x_n}\}$ , the covariance matrix is:**

$$\Sigma = \begin{pmatrix} K_{\theta}(x_1, x_1) & K_{\theta}(x_1, x_2) & \cdots & K_{\theta}(x_1, x_n) \\ K_{\theta}(x_2, x_1) & K_{\theta}(x_2, x_2) & \cdots & K_{\theta}(x_2, x_n) \\ \vdots & \vdots & \vdots & \vdots \\ K_{\theta}(x_n, x_1) & K_{\theta}(x_n, x_2) & \cdots & K_{\theta}(x_n, x_n) \end{pmatrix}$$

...Which we can compute based on the input (and the parameters) alone!

# The Whole Picture

In practice, at **training time**

- Pick a **parameterized** kernel function  $K_{\theta}(x_i, x_j)$
- Collect training observations  $\bar{\mathbf{y}}_{\bar{\mathbf{x}}}$
- Optimize the kernel for maximum likelihood (e.g. via gradient descent)

Both the parameters  $\theta$  **and the observations**  $\bar{\mathbf{y}}_{\bar{\mathbf{x}}}$  are stored in the model

- This is similar to what we have in Kernel Density Estimation

# The Whole Picture

## In practice, at **training time**

- Pick a **parameterized** kernel function  $K_{\theta}(x_i, x_j)$
- Collect training observations  $\bar{y}_{\bar{x}}$
- Optimize the kernel for maximum likelihood (e.g. via gradient descent)

Both the parameters  $\theta$  **and the observations**  $\bar{y}_{\bar{x}}$  are stored in the model

- This is similar to what we have in Kernel Density Estimation

## At inference time:

- Given a new input (i.e. index) value  $x$
- We obtain the covariance matrix  $\Sigma_{\bar{x}}$
- We obtain the covariance matrix  $\Sigma_{x, \bar{x}}$


...And with this we can completely characterize  $f(y_x, | \bar{y}_{\bar{x}})$

# Gaussian Processes in scikit-learn

## Let's see how to use Gaussian Processes in scikit-learn

First, let us choose a function as our ground truth

```
In [2]: f = lambda x: x * np.sin(2*np.pi*x) + x # target function  
x = np.linspace(0, 3, 1000)  
y = pd.Series(index=x, data=f(x))  
util.plot_gp(target=y, figsize=figsize)
```






# Gaussian Processes in scikit-learn

Then we build a small training set:

```
In [3]: np.random.seed(42)
n_tr = 15
x_tr = np.linspace(0.2, 2.8, n_tr) + 0.2*np.random.rand(n_tr)
x_tr.sort()
y_tr = pd.Series(index=x_tr, data=f(x_tr) + 0.2*np.random.rand(n_tr))
util.plot_gp(target=y, samples=y_tr, figsize=figsize)
```



# Gaussian Processes in scikit-learn

Now we are where we would normally start in a data analysis

...And we need to choose a kernel among are many available options

- Since we have 15 training points
- This will be used at training time to build a  $15 \times 15$  covariance matrix

We will start with a simple **Radial Basis Function** (i.e. Gaussian) kernel

$$K(x_i, x_j) = e^{-\frac{d(x_i, x_j)^2}{2l}}$$

The covariance decreases with the (Euclidean) distance  $d(x_i, x_j)$ :

- Intuitively, the closer the points, the higher the correlation
- The  $l$  parameter (scale) control the rate of the reduction

# Gaussian Processes in scikit-learn

## Here's how to use an RBF kernel in scikit-learn

```
In [4]: from sklearn.gaussian_process.kernels import RBF  
        kernel = RBF(1, (1e-2, 1e2))
```

**The RBF kernel has a single parameter, representing its **scale****

The extra (tuple) parameter represents a pair of **bounds**

- During training, only values within the boundaries will be considered

**Bounds can be very useful for controlling the training process**

- They should be chosen based on the available domain information
- Multiple educated guesses are fine when little is known about the problem

# Gaussian Processes in scikit-learn

## Now we can train a Gaussian Process

```
In [5]: from sklearn.gaussian_process import GaussianProcessRegressor
gp = GaussianProcessRegressor(kernel=kernel, n_restarts_optimizer=9)
gp.fit(y_tr.index.values.reshape(-1,1), y_tr.values) # needs 2D input
```

```
Out[5]: GaussianProcessRegressor
GaussianProcessRegressor(kernel=RBF(length_scale=1), n_restarts_optimizer=9)
```

- Training uses Gradient Descent to maximize the likelihood of the training data
- **Restarts** are needed to mitigate issues due to local optima

### Training can be an expensive process:

- Building the covariance matrix runs in  $O(n^2)$  (where  $n$  is the training set size)
- ...And we need to do that repeatedly during optimization!

# Gaussian Processes in scikit-learn

Finally, we can obtain the predictions:

```
In [6]: xp, std = gp.predict(x.reshape(-1,1), return_std=True)
xp = pd.Series(index=y.index, data=xp)
std = pd.Series(index=y.index, data=std)
```

The "predictions" are **not point estimates**!

- We actually get those (i.e. the means, called `xp` in the code)
- ...But we also obtain standard deviations

**The model output is a fully characterized conditional distribution**


The distribution is conditional on:

- The input value  $\mathbf{x}$  (passed at inference time)
- All the training observations (stored in the model)

# Gaussian Processes in scikit-learn

We can now plot the predictions

```
In [7]: util.plot_gp(target=y, samples=y_tr, pred=yp, std=std, figsize=figsize)
```



**Obtaining a full distribution as output is cool**

- We obtain an estimated value (the mean)
- ...And we get to know how certain the model is about that

## Behind the Scenes

### Let's see how prediction works with a numeric example

We want a prediction for  $x = 2.5$ , i.e.  $y_{2.5}$

- We consider two separate training set (and pre-optimized kernel parameters)
- The first set contains only  $(\bar{x}_9, \bar{y}_{\bar{x}_9}) \simeq (2.01, 2.27)$
- The second set contains only  $(\bar{x}_0, \bar{y}_{\bar{x}_0}) \simeq (0.27, 0.58)$

**The covariance matrices in the two cases are therefore:**

$$\Sigma_{y_x, \bar{y}_{\bar{x}_9}} = \begin{pmatrix} K(2.01, 2.01) & K(2.01, 2.5) \\ K(2.5, 2.01) & K(2.5, 2.5) \end{pmatrix}$$

$$\Sigma_{y_x, \bar{y}_{\bar{x}_0}} = \begin{pmatrix} K(0.27, 0.27) & K(0.27, 2.5) \\ K(2.5, 0.27) & K(2.5, 2.5) \end{pmatrix}$$

# Behind the Scenes

## Let's actually build the matrices in Python

- Note: scikit-learn kernels are not designed to be used on individual points
- So, for this we will rely on basic numpy methods

We start with  $\bar{x}_9$  and  $x$ , which are **close to each other**

```
In [9]: from scipy.stats import multivariate_normal
X9, X0, X = [[x_tr[9]]], [[x_tr[0]]], [[2.5]] # Must be 2D
sigma_9x = np.array([[kernel(X9, X9)[0,0], kernel(X9, X)[0,0]],
                    [kernel(X, X9)[0,0], kernel(X, X)[0,0]]])
f_9x = multivariate_normal([0, 0], cov=sigma_9x)
```

Then we do the same for  $\bar{x}_0$  and  $x$ , which are **far apart**


```
In [10]: sigma_0x = np.array([[kernel(X0, X0)[0,0], kernel(X0, X)[0,0]],
                             [kernel(X, X0)[0,0], kernel(X, X)[0,0]]])
f_0x = multivariate_normal([0, 0], cov=sigma_0x)
```



## Behind the Scenes

$\bar{x}_9$  and  $x$  are **close to each other**, so  $\bar{y}_{\bar{x}_9}$  and  $y_x$  are **strongly correlated**

```
In [11]: yr = np.linspace(-5, 5, 100)
util.plot_distribution_2D(f_9x, yr, yr, figsize=figsize)
plt.xlabel('y_{x_9}'); plt.ylabel('y_{x}'); plt.tight_layout()
```




- If we know **neither**  $\bar{y}_{\bar{x}_9}$  nor  $y_x$ , we can only say that they are likely **both zero**

# Behind the Scenes

But we **do know**  $\bar{y}_{\bar{y}_9}$  ! So, we can use this information

```
In [12]: util.plot_distribution_2D(f_9x, yr, yr, figsize=figsize)
plt.axvline(10*(y_tr[x_tr[9]] + 5), color='tab:orange');
plt.xlabel('y_{x_9}'); plt.ylabel('y_{x}'); plt.tight_layout()
```




- Given the observation, the most likely value for  $y_x$  is  $\simeq 2$

## Behind the Scenes

$\bar{x}_0$  and  $x$  are **far apart**, so  $\bar{y}_{\bar{x}_0}$  and  $y_x$  are **loosely correlated**

```
In [13]: util.plot_distribution_2D(f_0x, yr, yr, figsize=figsize)
plt.axvline(10*(y_tr[x_tr[0]] + 5), color='tab:orange');
plt.xlabel('y_{x_0}'); plt.ylabel('y_{x}'); plt.tight_layout()
```



- Knowing  $\bar{y}_{\bar{x}_0}$  is not going to be of much help here

# Memento

**So, a few key insight to keep in mind:**

- Superficially, GPs behave like functions that output probability distribution
- Internally, they store the kernel, defining how all the points are correlated
- A set of observations, used to obtain conditional distributions

**In scikit-learn:**

When we call the `fit` method:

- The optimizer adjusts the kernel parameters
- ...And the observations  $\hat{y}_x$  are stored

When we call the `predict` method:


- The covariance matrix is built
- The model computes the conditional distributions

# How to Improve the Model

**We can improve the model by choosing a more appropriate kernel**

- Let's (cheat and) look again at our ground truth and training data
- Which patterns and properties do you notice?

```
In [14]: util.plot_gp(target=y, samples=y_tr, figsize=figsize)
```




# How to Improve the Model

**We can improve the model by choosing a more appropriate kernel**

- Let's (cheat and) look again at our ground truth and training data
- Which patterns and properties do you notice?

```
In [14]: util.plot_gp(target=y, samples=y_tr, figsize=figsize)
```



**We have some noise, a period, and a trend**

# How to Improve the Model

So, let us deal with the noise first

```
In [15]: from sklearn.gaussian_process.kernels import WhiteKernel  
  
kernel = WhiteKernel(0.1, (1e-2, 1e2))  
kernel += RBF(1, (1e-2, 1e2))
```

**WhiteKernel** captures the presence of **noise** in the data

$$K(x_i, x_j) = \sigma^2 \text{ iff } x_i = x_j, 0 \text{ otherwise}$$

- The only parameter of `WhiteKernel` represents the noise level  $\sigma^2$
- A small noise level prevents overfitting
- ...But too much noise leads to useless predictions!

# How to Improve the Model

It's often a good idea to have **magnitude parameters** in the kernel

```
In [16]: from sklearn.gaussian_process.kernels import ConstantKernel  
  
kernel = WhiteKernel(0.1, (1e-2, 1e2))  
kernel += ConstantKernel(1, (1e-2, 1e2)) * RBF(1, (1e-2, 1e2))
```

**ConstantKernel** is a constant factor (in this case a relative weight)

- ...And allows the optimizer to tune the magnitude of the RBF kernel



# How to Improve the Model

It's often a good idea to have **magnitude parameters** in the kernel

```
In [16]: from sklearn.gaussian_process.kernels import ConstantKernel  
  
kernel = WhiteKernel(0.1, (1e-2, 1e2))  
kernel += ConstantKernel(1, (1e-2, 1e2)) * RBF(1, (1e-2, 1e2))
```

**ConstantKernel** is a constant factor (in this case a relative weight)

■ ...And allows the optimizer to tune the magnitude of the RBF kernel


**Let's repeat training again:**

```
In [17]: gp = GaussianProcessRegressor(kernel=kernel, n_restarts_optimizer=9)  
gp.fit(y_tr.index.values.reshape(-1,1), y_tr.values) # needs 2D input  
print(gp.kernel_)  
  
WhiteKernel(noise_level=0.01) + 2.21**2 * RBF(length_scale=0.321)
```

# How to Improve the Model

## Let us see the new predictions

```
In [18]: xp, std = gp.predict(x.reshape(-1,1), return_std=True)
util.plot_gp(target=y, samples=y_tr, pred=pd.Series(index=y.index, data=xp),
             std=pd.Series(index=y.index, data=std), figsize=figsize)
```



- Better, since the black curve is mostly in the confidence interval
- ...But we are still not exploiting the period and the trend

# How to Improve the Model

So, let us take them into account, starting with the **period**

```
In [19]: from sklearn.gaussian_process.kernels import ExpSineSquared
kernel = WhiteKernel(0.1, (1e-2, 1e2))
kernel += ConstantKernel(1, (1e-2, 1e2)) * RBF(1, (1e-2, 1e2))
kernel += ExpSineSquared(1, 1, (1e-1, 1e1), (1e-1, 1e1))
```

**ExpSineSquared** captures the period:

$$K(x_i, x_j) = e^{-2 \frac{\sin^2\left(\pi \frac{d(x_i, x_j)}{p}\right)}{l^2}}$$

- The correlation grows if the distance is close to a multiple of the period  $p$
- The scale parameter  $l$  controls the rate of decrease/increase
- In the implementation, the first parameter is  $l$  and the second  $p$

# How to Improve the Model

Now, let's try to capture the **trend**

```
In [20]: from sklearn.gaussian_process.kernels import DotProduct
kernel = WhiteKernel(0.1, (1e-2, 1e2))
kernel += ConstantKernel(1, (1e-2, 1e2)) * RBF(1, (1e-2, 1e2))
kernel += ExpSineSquared(1, 1, (1e-1, 1e1), (1e-1, 1e1))
kernel += DotProduct(1, (1e-2, 1e2))
```

**DotProduct (somewhat) captures the trend:**


$$K(x_i, x_j) = \sigma^2 + x_i x_j$$

- The larger the  $x$  values, the larger the correlation
- This allows the distance from the mean (which is zero) to grow
- The  $\sigma$  parameter controls the base level of correlation
- Unlike all kernels so far DotProduct is **not translation-invariant**

# How to Improve the Model

The new predictions are a bit better at the edges of the plot

```
In [21]: gp = GaussianProcessRegressor(kernel=kernel, n_restarts_optimizer=9)
gp.fit(y_tr.index.values.reshape(-1,1), y_tr.values) # needs 2D input
xp, std = gp.predict(x.reshape(-1,1), return_std=True)
util.plot_gp(target=y, samples=y_tr, pred=pd.Series(index=y.index, data=xp),
             std=pd.Series(index=y.index, data=std), figsize=figsize)
```



- Notice how the C.I. grows close to the sides?
- GPs tend to **interpolate well** and **extrapolate poorly**